

On bayesian binary probit regression

Feb 16 2020

R Markdown

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com> (<http://rmarkdown.rstudio.com>).

Model set-up:

Let $\mathbf{y} = (y_1, y_2, \dots, y_n)^T \in \{0, 1\}$, and we try to model $P(Y_i = 1) = \Phi(\mathbf{x}_i^T \boldsymbol{\beta})$

Introduce n latent variables $\mathbf{z} = (z_1, z_2, \dots, z_n)^T$, and that $z_i = \mathbf{x}_i^T \boldsymbol{\beta} + \epsilon_i$, $\epsilon_i \stackrel{iid}{\sim} N(0, 1^2)$

Then we have:

$$Y_i = \begin{cases} 1, & \text{if } z_i > 0 \\ 0, & \text{if } z_i \leq 0 \end{cases}$$

since $P(Y_i = 1) = P(z_i > 0) = P(\mathbf{x}_i^T \boldsymbol{\beta} + \epsilon_i > 0)$

$$= P(\epsilon_i > -\mathbf{x}_i^T \boldsymbol{\beta}) = 1 - P(\epsilon_i \leq -\mathbf{x}_i^T \boldsymbol{\beta}) = P(\epsilon_i \leq \mathbf{x}_i^T \boldsymbol{\beta}) = \Phi(\mathbf{x}_i^T \boldsymbol{\beta})$$

Thus, the unknown parameters of interest are $(\boldsymbol{\beta}, \mathbf{Z})$, and

$$y_i \sim \text{bernoulli}(\Phi(\mathbf{x}_i^T \boldsymbol{\beta})) = \text{bernoulli}(P(z_i > 0)), \mathbf{Z} \sim N_n(\mathbf{x}_i^T \boldsymbol{\beta}, I_n), \boldsymbol{\beta} \sim N_p(\boldsymbol{\beta}_0, \Sigma_0).$$

The joint posterior density of $(\boldsymbol{\beta}, \mathbf{Z})$ is:

$$\begin{aligned} p(\boldsymbol{\beta}, \mathbf{Z} | \mathbf{y}, X) &\propto p(\boldsymbol{\beta}, \mathbf{Z}, \mathbf{y} | X) = p(\mathbf{y} | \mathbf{Z}, \boldsymbol{\beta}, X) p(\mathbf{Z} | \boldsymbol{\beta}, X) p(\boldsymbol{\beta} | X) \\ &= p(\mathbf{y} | \mathbf{Z}) p(\mathbf{Z} | \boldsymbol{\beta}, X) p(\boldsymbol{\beta}) \end{aligned}$$

$$= \prod_{i=1}^n \{P(z_i > 0)^{y_i} [1 - P(z_i > 0)]^{1-y_i}\} |2\pi I_n|^{-\frac{1}{2}} \exp(-\frac{1}{2}(\mathbf{z} - X\boldsymbol{\beta})^T (\mathbf{z} - X\boldsymbol{\beta})) |2\pi \Sigma_0|^{-\frac{1}{2}} \exp(-\frac{1}{2}(\boldsymbol{\beta} - \boldsymbol{\beta}_0)^T \Sigma_0^{-1} (\boldsymbol{\beta} - \boldsymbol{\beta}_0)) - (1)$$

From (1) we can derive the full conditionals for $[\boldsymbol{\beta} | \mathbf{z}, X, \mathbf{y}]$, $[\mathbf{Z} | \boldsymbol{\beta}, X, \mathbf{y}]$:

$$\begin{aligned} p(\boldsymbol{\beta} | \mathbf{z}, X, \mathbf{y}) &\propto p(\boldsymbol{\beta}, \mathbf{Z}, \mathbf{y} | X) \\ &\propto \exp(-\frac{1}{2}[(\boldsymbol{\beta}^T X^T - \mathbf{z}^T)(X\boldsymbol{\beta} - \mathbf{z}) + (\boldsymbol{\beta}^T - \boldsymbol{\beta}_0^T) \Sigma_0^{-1} (\boldsymbol{\beta} - \boldsymbol{\beta}_0)]) \\ &= \exp(-\frac{1}{2}[\boldsymbol{\beta}^T (\Sigma_0^{-1} + X^T X) \boldsymbol{\beta} - 2\boldsymbol{\beta}^T (\Sigma_0^{-1} \boldsymbol{\beta}_0 + X^T \mathbf{z}) + \mathbf{z}^T \mathbf{z} + \boldsymbol{\beta}_0^T \Sigma_0^{-1} \boldsymbol{\beta}_0]) \end{aligned}$$

$$\text{Let } Q_{\boldsymbol{\beta}} = \Sigma_0^{-1} + X^T X, l_{\boldsymbol{\beta}} = \Sigma_0^{-1} \boldsymbol{\beta}_0 + X^T \mathbf{z},$$

it can be recognized that $[\boldsymbol{\beta} | \mathbf{z}, X, \mathbf{y}] \sim N_p(Q_{\boldsymbol{\beta}}^{-1} l_{\boldsymbol{\beta}}, Q_{\boldsymbol{\beta}}^{-1})$

$$\begin{aligned} p(\mathbf{Z} | \boldsymbol{\beta}, X, \mathbf{y}) &\propto p(\boldsymbol{\beta}, \mathbf{Z}, \mathbf{y} | X) \\ &\propto p(\mathbf{y} | \mathbf{z}) p(\mathbf{z} | \boldsymbol{\beta}, X) \\ &\propto \prod_{i=1}^n \{P(z_i > 0)^{y_i} [1 - P(z_i > 0)]^{1-y_i}\} \exp(-\frac{1}{2}(\mathbf{z} - X\boldsymbol{\beta})^T (\mathbf{z} - X\boldsymbol{\beta})) \end{aligned}$$

where $p(y_i | z_i) = I(y_i = 1)I(z_i > 0) + I(y_i = 0)I(z_i \leq 0)$

$$[z_i | \boldsymbol{\beta}, y_i, X] \sim \begin{cases} N(z_i \in (0, +\infty); \mathbf{x}_i^T \boldsymbol{\beta}, 1), & \text{if } y_i = 1 \\ N(z_i \in (-\infty, 0); \mathbf{x}_i^T \boldsymbol{\beta}, 1), & \text{if } y_i = 0 \end{cases}$$

Implementation

```

set.seed(20740)
# sample size
n= 100
# dim of reg. parameters
p= 2
x= runif(n= n, min= 0, max= 1)

xb= -2+ 6*x

calcProb= function(x) {
  return(exp(xb)/(1+ exp(xb)))
}
pi.x= calcProb(x= xb)

set.seed(21218)
y= rbinom(n= n, size= 1, prob= pi.x)

fit= glm(y ~ x, family= "binomial")
summary(fit)

```

```

##
## Call:
## glm(formula = y ~ x, family = "binomial")
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.3508  -0.6465   0.2379   0.5827   2.1486
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -2.3435     0.5653  -4.145 3.39e-05 ***
## x              6.7426     1.3273   5.080 3.78e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 131.791  on 99  degrees of freedom
## Residual deviance:  83.121  on 98  degrees of freedom
## AIC: 87.121
##
## Number of Fisher Scoring iterations: 5

```

```

n1= sum(y) # Number of successes
n0= n - n1 # Number of failures

X= matrix(c(rep(1, n), x), ncol= p)

require("mvtnorm")

```

```

## Loading required package: mvtnorm

```

```

require("truncnorm")

```

```

## Loading required package: truncnorm

```

```

# Conjugate prior on the coefficients beta ~ N(beta_0, Q_0)
beta.0= rep(0, p)
Q.0= diag(10, p)

# Initialize parameters
beta= rep(0, p)
z= rep(0, n)

# Number of simulations for Gibbs sampler
S= 15000
# Burn in period
nBurnin= 5000
# Matrix storing samples of the \theta parameter
post.beta= matrix(NA, nrow = S, ncol = p)

# -----
# Gibbs sampling algorithm
# -----
# Compute posterior variance of beta
# X'X, crossprod(X, X)= t(X)%*% X
Q.beta= t(X)%*% X+ solve(Q.0)
Q.beta.inv= solve(Q.beta)

set.seed(21218)
for (s in 1: S) {
  # Update Mean of z
  mu_z= X %*% beta
  # Draw latent variable z from its full conditional: z | \beta, y, X
  z[y == 0]= rtruncnorm(n0, mean = mu_z[y == 0], sd = 1, a = -Inf, b = 0)
  z[y == 1]= rtruncnorm(n1, mean = mu_z[y == 1], sd = 1, a = 0, b = Inf)

  # Compute posterior mean of theta
  M = Q.beta.inv %*% (t(X) %*% z+ solve(Q.0)%*% beta.0)
  # Draw variable \theta from its full conditional: \theta | z, X
  beta= c(rmvt(n= 1, mean= M, sigma= Q.beta.inv))

  # Store the \theta draws
  post.beta[s, ] = beta
}

# -----
# Get posterior mean of \theta
# -----
postMean= colMeans(post.beta[-(1: nBurnin), ]); postMean

```

```
## [1] -1.293848  3.761639
```

```

fit= glm(y ~ x, family = binomial(link = probit))
summary(fit)

```

```
##
## Call:
## glm(formula = y ~ x, family = binomial(link = probit))
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.3627  -0.6685   0.2062   0.6057   2.1345
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -1.3478     0.3093  -4.357 1.32e-05 ***
## x              3.8669     0.6851   5.645 1.66e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 131.791  on 99  degrees of freedom
## Residual deviance:  83.097  on 98  degrees of freedom
## AIC: 87.097
##
## Number of Fisher Scoring iterations: 5
```

```
mle.beta= fit$coefficients; mle.beta
```

```
## (Intercept)          x
##    -1.34779     3.86694
```

Note that the `echo = FALSE` parameter was added to the code chunk to prevent printing of the R code that generated the plot.